

# PRPROGRAMMING ASSIGNMENT 1

## Benchmarking

---

### 1. Introduction:

The assignment aims at Benchmarking different components of a computer system and thus, learn about the performance.

1. CPU Benchmarking
2. Memory Benchmarking
3. Disk Benchmarking

#### 1.1 CPU Benchmark:

1. CPU Benchmarking is done in C Language. It is also menu driven.
2. Goal of this module is to calculate processor speed in terms of Floating Point Operations Per Second (FLOPS) and Integer Operations Per Second (IOPS). The program runs multiple instructions concurrently.
3. Two functions namely flopsCal and iopsCal calculate the GFLOPS and GIOPS respectively. Time taken to perform operations is calculated using clock() function. The time difference thus obtained is divided by CLOCKS\_PER\_CYCLE giving time in seconds. Number of (ITERATIONS \* No of threads \* No of operations) are divided by calculated time to determine processor speed in terms of Giga FLOPS and Giga IOPS. Functions threadFunctionFlopCal and threadFunctionIopsCal are used to create and join threads. For the purpose of multi-threading pthread libraries are used. The loop iterating variables are volatile to prevent the compiler optimizing the loops. It is difficult to calculate CPU speed in single iteration, hence taken 1000000000.
4. It also includes the experiment to calculate the GFLOPS/GIOPS over a period of 10 mins for 4 threads. The functions flops and iops calculate the instructions per second and store them in CPU\_LOG.txt and CPU\_IOPS.txt respectively. Thread functions like threadFunctionFlop and threadFunctionIops create, join and kill threads.
5. Improvements and extension to the program:  
Improve the overclocking ie. The processor to run faster than expected by giving more work for the CPU. Also, carry out experiments on increasing thread concurrency to give 100% utilization.

#### 1.2 Memory Benchmark:

1. Memory benchmarking program is developed in C language. It is also menu driven.
2. The two operations performed on memory are read and write. Also, read and write operations are performed sequentially and randomly. A random offset is used to perform read and write operations.
3. Memory is allocated using the malloc function.

4. Three block sizes (1 Byte, 1 KB, 1 MB) and concurrency (1 thread/2 thread) is considered.
5. Throughput in terms of MB/sec and latency in terms of milliseconds are calculated for 3 different block sizes of 1 Byte, 1 Kilobyte and 1 Megabyte at varying level of concurrency for 1 thread and 2 threads. Threading is implemented using `p_thread()` and `p_join()` functions of C language.
6. `clock()` function is used to determine latency. Difference of start time and end time is divided by `CLOCKS_PER_CYCLE` to calculate latency in seconds.
7. Read and Write is performed using C function `memcpy()`.
8. Pointer functions `*block_Byte()`, `*block_Kbyte()` and `*block_Mbyte()` are implemented to perform sequential data transfer of 1 Byte, 1 Kilobyte and 1 Megabyte respectively. Also, `*block_Byte_random()`, `*block_KByte_random()` and `*block_MByte_random()` functions are implemented to perform random data transfer of 1 Byte, 1 KB and 1 MB respectively.
9. Improvements and extension to the program:  
To test the read and write operations using functions other than `memcpy()`.  
Improve the code to isolate the effect of cache.

### 1.3 Disk Benchmark:

1. Disk Benchmarking program is developed in C language.
2. The benchmark aims at measuring the disk speed by performing read and write operations.
3. The functions `p_thread()` and `p_join()` are used to create and join multiple threads.
4. `clock()` function is used to determine latency. Difference of start time and end time is divided by `CLOCKS_PER_CYCLE` to calculate latency in seconds.
5. Throughput in terms of MB/sec and latency in terms of milliseconds are calculated for 3 different block sizes of 1 Byte, 1 Kilobyte and 1 Megabyte at varying level of concurrency for 1, 2 and 4 threads respectively.
6. Read and Write is performed on data of block size 1 Byte, 1 KB and 1MB using C function `fseek`, `fputc`, `fread()` and `fwrite()` respectively.
7. `fseek` sets the indicator to a new position either at beginning of the file or at any random position in the file.
8. Improvements and extension to the program:  
Experiment can be carried out by increasing the concurrency and minimizing the latency as far as possible. Lines of code can be optimized.

## 2. Performance


Performance for various modules is measured and plotted on graphs.

### 2.1 CPU Benchmarking Performance:

1. The floating point operation is in the range of 7-9 GFLOPS for a single thread.
2. The integer point operation is in the range of 9-10 GIOPS for a single thread.
3. Readings

THREADS	GFLOPS	GIOPS
1	10.308567	10.304895
2	10.293027	10.294354
3	10.268674	10.28658

Graph 1: GFLOPS/GIOPS vs No of threads



Observation 1: The GFLOPS and GIOPS decreases as the number of threads increase.

Observation 2: It is observed that GLOPS and GIOPS are maximum when run on single thread.

Observation 3: Optimal performance is obtained at 1 thread

4. Calculating theoretical peak performance
  - a. T2 micro instance Specification
    - CPU speed: 2.6 GHz
    - Number of Cores: 1
    - No. of instruction per cycle: 16 (Assumption)

$$\begin{aligned}\text{Theoretical Peak Performance} &= (\text{CPU Speed} * \text{No of cores} * \text{Instructions per cycle}) \\ &= 41.6 \text{ GFLOPS}\end{aligned}$$

$$\begin{aligned}\text{Efficiency} &= (\text{FLOPS for 1 thread} / \text{Theoretical Peak Performance}) * 100 \\ &= (10.308567 / 41.6) * 100 \\ &= 24.7 \%\end{aligned}$$

b. Computer Specifications

Processor: 4X Intel(R) core(TM) i3 CPU.

CPU speed: 2.40 GHz

Number of Cores: 2


No. of instruction per cycle: 4

$$\begin{aligned}\text{Theoretical Peak Performance} &= (\text{CPU Speed} * \text{No of cores} * \text{Instructions per cycle}) \\ &= 19.2 \text{ GFLOPS}\end{aligned}$$

$$\begin{aligned}\text{Efficiency} &= (\text{FLOPS for 1 thread} / \text{Theoretical Peak Performance}) * 100 \\ &= (8.9 / 19.2) * 100 \\ &= 46.35 \%\end{aligned}$$

5. Plot: GFLOPS over a period of 10 mins with four threads running simultaneously.

Graph 2: GFLOPS sec vs Time in seconds




Observation 1: The GFLOPS are in the range of 7 to 93

Observation 2: At 511<sup>th</sup> second the GFLOPS are maximum.

6. Plot: GIOPS over a period of 10 mins with four threads running simultaneously.

Graph 3: Instructions per sec vs Time in seconds



Observation 1: The GIOPS are the range of 2 – 48.

Observation 2: At 369<sup>th</sup> second the GIOPS are maximum.

### Linpack Benchmark:

The following are the observations obtained when LINPACK is run on AMAZON AWS t2 micro instance.

```

ubuntu@ip-172-31-62-7: /usr$
Number of cores: 1
Number of threads: 1

Parameters are set to:
Number of tests: 15
Number of equations to solve (problem size) : 1000 2000 5000 10000 15000 18000 20000 22000 25000 26000 27000 30000 35000 40000 45000
Leading dimension of array : 1000 2000 5000 10000 15000 18000 20016 22008 25000 26000 27000 30000 35000 40000 45000
Number of trials to run : 4 2 2 2 2 2 2 2 2 2 1 1 1 1 1
Data alignment value (in Kbytes) : 4 4 4 4 4 4 4 4 4 4 4 1 1 1 1

Maximum memory requested that can be used=800204096, at the size=10000

===== Timing linear equation system solver =====
Size LDA Align. Time(s) GFlops Residual Residual(norm) Check
1000 1000 4 0.025 27.2373 7.441825e-13 2.537853e-02 pass
1000 1000 4 0.023 29.2076 7.441825e-13 2.537853e-02 pass
1000 1000 4 0.023 28.9412 7.441825e-13 2.537853e-02 pass
1000 1000 4 0.023 29.0702 7.441825e-13 2.537853e-02 pass
2000 2000 4 0.170 31.4176 3.616191e-12 3.145643e-02 pass
2000 2000 4 0.170 31.3603 3.616191e-12 3.145643e-02 pass
5000 5000 4 2.470 33.7643 2.067851e-11 2.883452e-02 pass
5000 5000 4 2.446 34.0904 2.067851e-11 2.883452e-02 pass
10000 10000 4 18.637 35.7825 6.859494e-11 2.418727e-02 pass
10000 10000 4 18.486 36.0732 6.859494e-11 2.418727e-02 pass

Performance Summary (GFlops)
Size LDA Align. Average Maximal
1000 1000 4 28.6141 29.2076
2000 2000 4 31.3889 31.4176
5000 5000 4 33.9273 34.0904
10000 10000 4 35.9278 36.0732

Residual checks PASSED

End of tests

Done: Fri Feb 12 21:06:38 UTC 2016
ubuntu@ip-172-31-62-7: /usr$

```

Average GFLOPS = 31.695

Best performance = 36.0732 GFLOPS

- a. Observation 1: Efficiency with respect to theoretical performance of t2 micro instance  
 Efficiency = (FLOPS for 1 thread / Theoretical Peak Performance) \* 100  
 = (31.695/41.6) \* 100  
 = 76.18 %
- b. Observation 2: Efficiency of my code with respect to linpack  
 Efficiency = (FLOPS for 1 thread / Theoretical Peak Performance) \* 100  
 = (10.3087/31.695) \* 100  
 = 32.52%

## 2.2 Memory Benchmarking Performance:

1. Readings for Sequential Read Write operations varying block sizes and varying concurrency.

Sequential Read Write	Thread 1	Thread 2
1byte	195.121951	193.89239
1Kb	3457.62712	2170.23585
1 Mb	2288.78276	3630.67

2. Readings for Random Read Write operations varying block sizes and varying concurrency.

Random Read Write	Thread 1	Thread 2
1byte	26.295722	26.053879
1Kb	1610.94177	1354.56368
1 Mb	16447.8887	11616.27

3. Theoretical Calculation

Processor: Intel(R) Core i3-2330M


Clock Speed: 100 MHz

Memory Type: DDR3-1066/1333

Theoretical Maximum Memory Bandwidth is calculated using formula:

Base DRAM clock frequency \* lines per clock \* memory bus width \* number of interfaces = 133000000 \* 2 \* 64 \* 2 = 4256 Mb/s

4. Graph 4: Display throughput for Sequential Read & Write

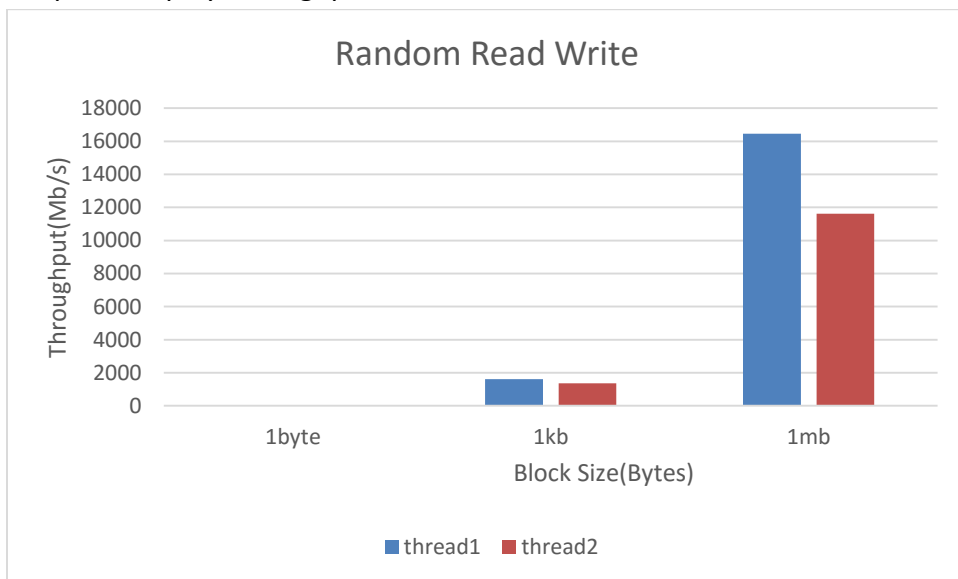


Observation:

It is observed that throughput increases with increase in size of block.

Throughput is higher for more number of threads.

5. Graph 5. Display throughput for Random Read & Write




Observation:

It is observed that throughput increases with increase in size of block.

Throughput is higher for single thread.


6. Graph 6: Display Latency for Sequential Read & Write



Observation:

It is observed that latency decrease with increase in block size. This is applicable for 1 & 2 threads. Latency is slightly higher for 2 threads.

7. Graph 7: Random Read write latency



Observation:

It is observed that latency decrease with increase in block size. This is applicable for 1 & 2 threads. Latency is slightly higher for 2 threads.




8. Graph 8: Compare Sequential Read Write to Random Read Write with respect to throughput



Observation:

It is observed that throughput increases with increase in block size. This is applicable for 1 & 2 threads. Throughput is higher for Rand Read Write.

9. Compare Sequential Read Write to Random Read Write with respect to latency



Observation:

It is observed that latency decreases with increase in block size. This is applicable for 1 & 2 threads. Latency is highest to Rand Read Write 1 Byte of data.

## Stream Benchmark:

```

ubuntu@ip-172-31-62-7: ~
ubuntu@ip-172-31-62-7:~$ ./stream
-----
STREAM version $Revision: 5.10 $
-----
This system uses 8 bytes per array element.
-----
Array size = 10000000 (elements), Offset = 0 (elements)
Memory per array = 76.3 MiB (= 0.1 GiB).
Total memory required = 228.9 MiB (= 0.2 GiB).
Each kernel will be executed 10 times.
The *best* time for each kernel (excluding the first iteration)
will be used to compute the reported bandwidth.
-----
Your clock granularity/precision appears to be 1 microseconds.
Each test below will take on the order of 28919 microseconds.
(= 28919 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----
WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----
Function      Best Rate MB/s  Avg time     Min time     Max time
Copy:         5596.2      0.028851     0.028591     0.029064
Scale:        5497.7      0.029438     0.029103     0.029663
Add:          7929.2      0.030875     0.030268     0.031186
Triad:        7256.9      0.033216     0.033072     0.033444
-----
Solution Validates: avg error less than 1.000000e-13 on all three arrays
-----
ubuntu@ip-172-31-62-7:~$

```

Best performance achieved is 5596.2 MB/s

The theoretical performance calculated for my system is 4256 Mb/s. But the stream benchmark is executed on Amazon t2 micro instance. The specifications must be higher than the ones achieved on my system. Stream has achieved 75% or more efficiency.

## 2.3 Disk Benchmarking Performance:

### 1. Readings

Random write	Thread 1	Thread 2
1byte	1.264756	1.24779
1Kb	821.686068	834.43091
1 Mb	3471.01701	3595.82884

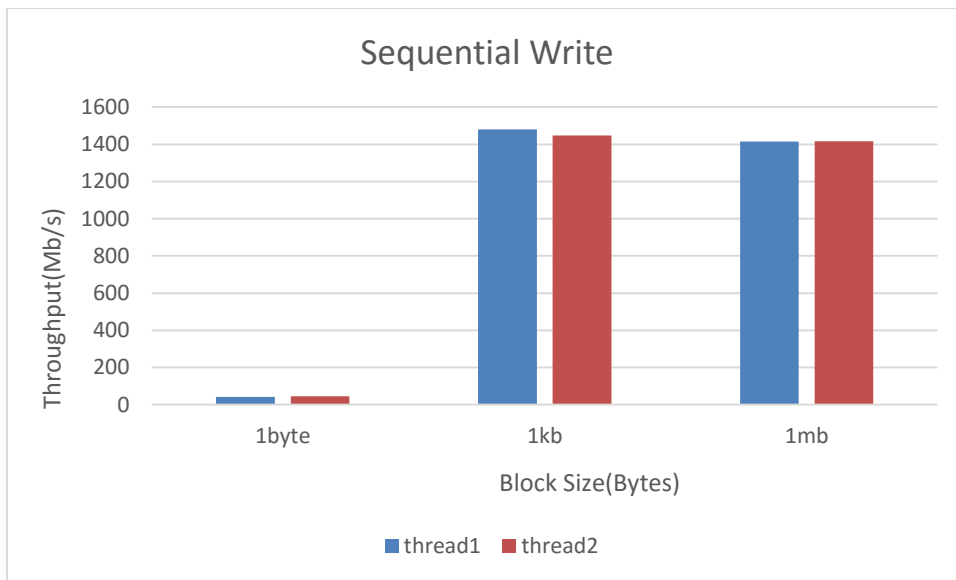
Sequential write	Thread 1	Thread 2
1byte	41.777726	45.702168
1Kb	1479.95531	1447.47579
1 Mb	1415.42817	1416.93234

Random read	Thread 1	Thread 2
1byte	6.178667	6.40438
1Kb	4770.14451	4775.95168
1 Mb	8928.57143	10422.0948

Sequential Read	Thread 1	Thread 2
1byte	33.288505	32.852159
1Kb	3837.03871	4287.06525
1 Mb	5649.71751	7446.01638


2. Sequential Write operation for varying block sizes, for varying concurrency (1 thread/ 2 thread). It is observed that the throughput increases with the block size. Also, the throughput values are almost same for 1 and 2 threads. Throughput is maximum for 1 kb.

Graph 9: Throughput vs Block Size for Sequential Write



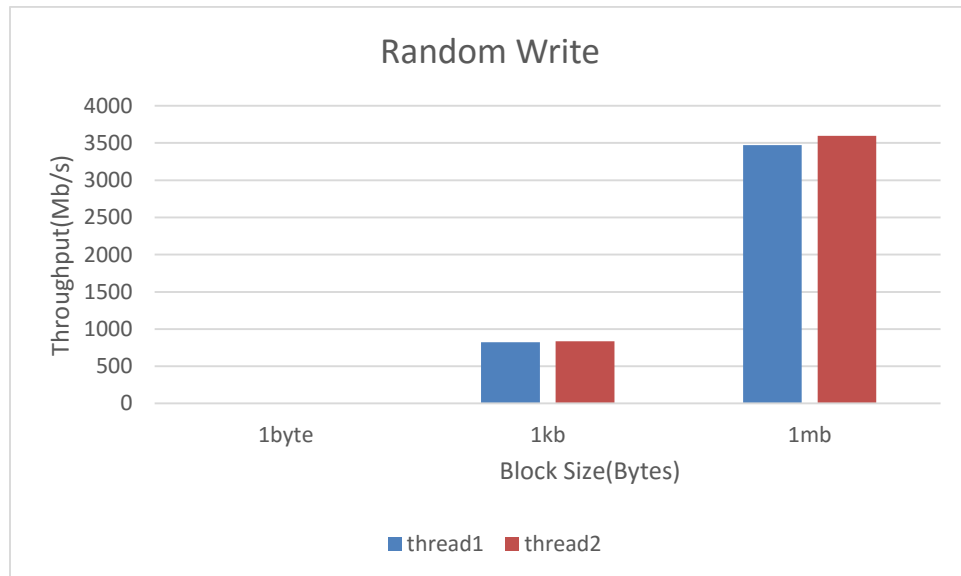
3. Sequential Write Operation for varying block size, for varying concurrency (1 thread, 2 thread). It is observed that the latency is maximum for 1 Byte and decreases with increase in block size. Also, the Latency is almost same for 1 and 2 threads.

Graph 10: Latency vs Block Size for Sequential Write



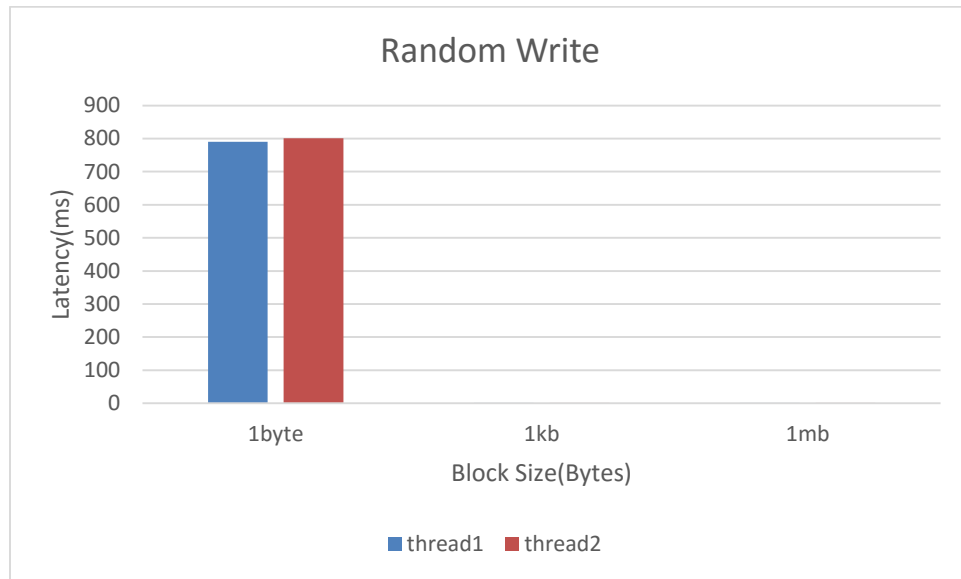
4. Random Write Operation for varying block size, for varying concurrency (1 thread, 2 thread). It is observed that the throughput is extremely high for 1 MB and decreases drastically with decrease in block size. Also, the throughput is almost same for 1 and 2 threads.

Graph 11: Throughput vs Block Size for Random Write




5. Random Write Operation for varying block size, for varying concurrency (1 thread, 2 thread). It is observed that the latency is maximum for 1 Byte and is almost negligible as the block size increases. Also, the Latency is almost same for 1 and 2 threads.

Graph 12: Latency vs Block Size for Random Write



6. Sequential Read Operation for varying block size, for varying concurrency (1 thread, 2 thread). It is observed that the throughput is extremely high for 1 MB and decreases drastically with decrease in block size. Also, the throughput is almost same for 1 and 2 threads. Throughput is maximum for thread no 2 with increase in block size.


Graph 8: Throughput vs Block Size for Sequential Read



7. Sequential Read Operation for varying block size, for varying concurrency (1 thread, 2 thread). It is observed that the latency is extremely high for 1 Byte and decreases


drastically with increase in block size. Also, the latency is almost same for 1 and 2 threads.

Graph 9: Latency vs Block Size for Sequential Write



- Random Read Operation for varying block size, for varying concurrency (1 thread, 2 thread). It is observed that the throughput is extremely high for 1 MB and decreases drastically with decrease in block size. Also, the throughput is almost same for 1 and 2 threads.


Graph 10: Throughput vs Block Size for Sequential Write



- Random Read Operation for varying block size, for varying concurrency (1 thread, 2 thread). It is observed that the latency is extremely high for 1 Byte and decreases

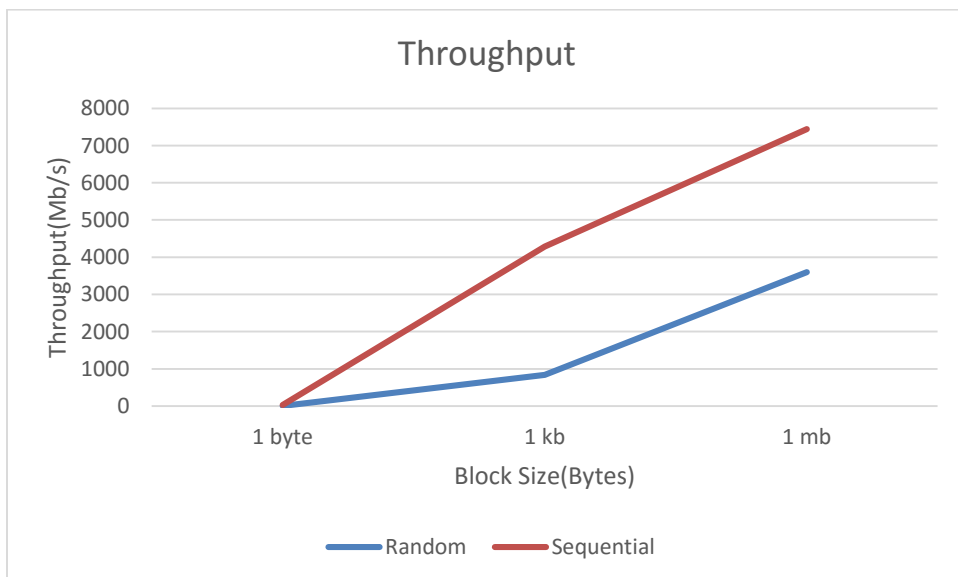
drastically with increase in block size. Also, the latency is almost same for 1 and 2 threads.

Graph 11: Latency vs Block Size for Sequential Write




10. Throughput for Sequential Access Read Operation is more than throughput for Random Access Read Operation for 1 thread. Same observed for 2 threads.

Graph 12: Throughput vs Block Size for Comparing Sequential and Random Access throughput.




11. Throughput for Random Access Write Operation is more than throughput for Sequential Access Write Operation for 1Mb for 1 thread. Similar observations are made for 2 threads.

Graph 13: Throughput vs Block Size for Comparing Sequential and Random Access throughput for Write operation 1 thread.



12. Latency for Random Access Write Operation is more than latency for Sequential Access Write Operation for 1 thread for 1 byte. Similar observation for 2 threads.


Graph 14: Latency vs Block Size for Comparing Sequential and Random Access throughput for Write operation for 1 thread. Similar observation for 2 threads.



13. Latency for Random Access Read Operation is more than latency for Sequential Access Read Operation for 1 thread. Similar observation for 2 threads.



Graph 15: Latency vs Block Size for Comparing Sequential and Random Access throughput for Write operation for 1 thread.



14. Optimal number of concurrency for best performance is 2.

15. More the number of threads more throughput.

### IOZone Benchmark:

- IOZone Benchmark is run for verifying outcomes of the implemented benchmark program

Operation	IOZone(Mb/s)	Observation(Mb/s)	Efficiency
Random Write	5710.61	828.055	14.5%
Random Read	10454.984	4773.04	45.65%
Sequential Read	10769.573	4062.0475	37.71%
Sequential Write	2458.618	1463.71	59.53%

```

ubuntu@ip-172-31-62-7: ~
16384 4096 2285711 3117688 8270680 1044876111522755 5200975 7089740 6166607 11545987 4096100 3521221 7223891 8613831
16384 8192^C
lozone: interrupted
Excel output is below:
"Writer report"
"4" "8" "16" "32" "64" "128" "256" "512" "1024" "2048" "4096" "8192" "16384"
"64" 1204796 1828508 1879725 1947927 1638743
"128" 1319723 1421023 2211113 2202044 2132084 2211113
"256" 1642294 2325094 2366082 2639446 2325094 2463808 2392442
"512" 1867692 2413429 2360376 2497638 2584820 2427068 2548017 2521095
"1024" 1939207 2365323 2138070 1815435 2540189 2617596 1825466 2567523 2414523
"2048" 1856959 2216868 2045804 2107025 2528082 1982533 2455104 2273187 2458618 2441150
"4096" 1754156 2149954 2294686 2207980 2183288 2137648 2256706 2290708 1991221 2249320 2264141
"8192" 1758718 2007837 2307601 1957060 2116815 2256085 2081297 2010304 2109796 2323361 2345886 2028223
"16384" 1726102 1954656 2117650 2283509 2366633 2268657 2548771 1971592 2379664 2406160 2285711 0
"Re-writer report"
"4" "8" "16" "32" "64" "128" "256" "512" "1024" "2048" "4096" "8192" "16384"
"64" 4018152 4018152 4897948 4897948 4274062
"128" 4012317 4759253 4934216 5122535 5545860 4717434
"256" 4264168 5022044 3454704 5347168 5569035 5217259 5242734
"512" 4262523 4784885 4784885 5453155 5624545 5822804 4742616 5067142
"1024" 4033569 4211554 4415030 5536137 4146499 5821271 5120336 5417427 5424269
"2048" 3121562 4016229 4014352 5507054 4751580 3899542 3307884 3513544 3312988 2990085
"4096" 2556982 2803592 2921346 2933818 3146031 3185110 3247731 3027937 3029528 2989984 2974454
"8192" 2434477 2711724 2926772 2988893 2918320 3100833 3025209 2942563 3031882 3011685 2975692 2915349
"16384" 2558831 2777464 2963220 3053818 3057214 3103607 3232822 3071562 3165365 3091877 3117688 0
"Reader report"
"4" "8" "16" "32" "64" "128" "256" "512" "1024" "2048" "4096" "8192" "16384"
"64" 12902017 15972885 15972885 15972885 15972885
"128" 11720614 9795896 14200794 14200794 10567140 15881078
"256" 11695808 14353744 14164395 13625180 15164624 15164624 13625180
"512" 11683444 14628067 14240069 14240069 15037801 13109944 10911694 9814569
"1024" 6009617 11645609 11368190 10429596 12983353 10990032 11132462 11368190 10769573
"2048" 8713618 9140117 13301113 10726615 8128206 6917293 7999524 8498106 9101380 12488897
"4096" 6245395 7380284 7473388 8028705 7983931 6883496 7274040 7418527 7874151 8531046 12264277
"8192" 6360787 6983962 7543623 7647724 7788135 7353127 7092077 7474698 7500806 8039611 8679319 11652086
"16384" 6384903 7291341 7468042 7568386 7710195 7282069 7059152 7467231 7278213 7820757 8270680 0

```

```

ubuntu@ip-172-31-62-7: ~
"Re-Reader report"
"4" "8" "16" "32" "64" "128" "256" "512" "1024" "2048" "4096" "8192" "16384"
"64" 12310336 22737791 20962191 15972885 20962191
"128" 5784891 15881078 15881078 14200794 16365173 15881078
"256" 12812277 16072608 15164624 14953435 17096249 15164624 13454450
"512" 12499490 15471149 14240069 14240069 16049269 14240069 12499490 12797441
"1024" 11398360 15080341 14422045 10990032 15080341 11773301 10990032 12348754 10878686
"2048" 11502234 15154989 12118884 15399510 14840791 13639023 12635867 12187663 9708272 12488897
"4096" 11099499 12151506 13658603 13790167 14524675 13086376 12125776 11477688 12117223 12643359 12841826
"8192" 10582628 8039611 12356037 13195968 9319649 10908526 11423523 10638329 10836280 11820445 11853067 11788003
"16384" 8928311 9799141 10937705 10220324 11251090 10421823 9571216 10495039 10246230 10965630 10448761 0
"Random read report"
"4" "8" "16" "32" "64" "128" "256" "512" "1024" "2048" "4096" "8192" "16384"
"64" 7940539 7940539 15972885 12902017 12902017
"128" 9129573 12542043 14200794 14200794 14200794 12542043
"256" 4496299 13454450 13454450 14164395 14953435 13454450 12812277
"512" 9510316 12797441 12797441 10434519 15037801 13872122 12215097 12797441
"1024" 9142007 12348754 11903823 13818817 14871477 13472053 11645609 12348754 10454984
"2048" 9149853 12729493 13574363 12416686 14017386 12416686 12118884 11900619 12050878 12327589
"4096" 9005041 12414943 12264277 13980948 14636041 13476460 12370246 12151506 10691939 12680688 12832234
"8192" 9329772 10694620 13298112 12717323 13934487 13065504 11977018 12409588 11656039 12048414 11667913 11098803
"16384" 8122096 10475840 13045180 12565719 14027802 12780725 11545987 11941233 11025448 10981401 11522755 0
"Random write report"
"4" "8" "16" "32" "64" "128" "256" "512" "1024" "2048" "4096" "8192" "16384"
"64" 3541098 4274062 4564786 4897948 4897948
"128" 2905056 4012317 5122535 5122535 5325799 5603747
"256" 4197489 5117791 5428265 5320671 5938650 5971678 5540300
"512" 4099771 4532414 5278892 5067142 5886650 5019764 5453155 5566231
"1024" 4029784 5169641 5690162 5277632 6025439 5782087 5821271 5536137 5885083
"2048" 4172290 4829046 5754221 6062888 5738844 5578583 5916724 5567735 5578583 5564129
"4096" 3846119 4858027 5127086 5891980 5826043 6094733 5777065 5625724 5841801 5627567 5452532
"8192" 4015955 4652131 5171695 5756551 5905961 5772993 5669170 5569014 5490700 5479317 5561803 5357988
"16384" 3610390 4157306 5099092 5278477 5883111 5824273 5862033 5643936 5561265 5557667 5200975 0
"Backward read report"
"4" "8" "16" "32" "64" "128" "256" "512" "1024" "2048" "4096" "8192" "16384"
"64" 6271021 9318832 3588436 7100397 6271021
"128" 6727225 8548124 6727225 7082197 9129573 8548124

```

